



+ Code + Text Copy to Drive

RAM Disk Editing

## Transfer Learning

In this notebook, you will perform transfer learning to train CIFAR-10 dataset on ResNet50 model available in Keras.

## Imports

```
[1] import os, re, time, json
    import PIL.Image, PIL.ImageFont, PIL.ImageDraw
    import numpy as np
try:
    # %tensorflow_version only exists in Colab.
    %tensorflow_version 2.x
except Exception:
    pass
import tensorflow as tf
from tensorflow.keras.applications.resnet50 import ResNet50
from matplotlib import pyplot as plt
import tensorflow_datasets as tfds

print("Tensorflow version " + tf.__version__)

Tensorflow version 2.5.0
```

## Parameters

- Define the batch size
- Define the class (category) names

```
[2] BATCH_SIZE = 32
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

Define some functions that will help you to create some visualizations. (These will be used later)

```
[3] #@title Visualization Utilities[RUN ME]
#Matplotlib config
plt.rcParams['image', cmaps='gray')
plt.rcParams['grid', linewidth=0]
plt.rcParams['xtick', top=False, bottom=True, labelsize='large')
plt.rcParams['ytick', left=False, right=True, labelsize='large')
plt.rcParams['axes', facecolor='F8F8F8', titlesize="large", edgecolor='white')
plt.rcParams['text', color='a8151a')
plt.rcParams['figure', facecolor="#F0F0F0"]# Matplotlib fonts
MPLTBLIB_FONT_DIR = os.path.join(os.path.dirname(__file__), "mpl-data/fonts/ttf")
# utility to display a row of digits with their predictions
def display_images(digits, predictions, labels, title):

    n = 10

    indexes = np.random.choice(len(predictions), size=n)
    n_digits = digits[indexes]
    n_predictions = predictions[indexes]
    n_predictions = n_predictions.reshape((n,))
    n_labels = labels[indexes]

    fig = plt.figure(figsize=(20, 4))
    plt.title(title)
    plt.yticks([])
    plt.xticks([])

    for i in range(10):
        ax = fig.add_subplot(1, 10, i+1)
        class_index = n_predictions[i]

        plt.xlabel(classes[class_index])
        plt.xticks([])
        plt.yticks([])
        plt.imshow(n_digits[i])

    # utility to display training and validation curves
    def plot_metrics(metric_name, title, ylim=5):
        plt.title(title)
        plt.ylim(0,ylim)
        plt.plot(history.history[metric_name], color='blue', label=metric_name)
        plt.plot(history.history['val_' + metric_name], color='green', label='val_' + metric_name)
```

## Visualization Utilities[RUN ME]

## Loading and Preprocessing Data

[CIFAR-10](#) dataset has 32 x 32 RGB images belonging to 10 classes. You will load the dataset from Keras.

```
[4] (training_images, training_labels), (validation_images, validation_labels) = tf.keras.datasets.cifar10.load_data()

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170500096/170498071 [=====] - 4s 0us/step
```

## Visualize Dataset

Use the `display_image` to view some of the images and their class labels.

```
[5] display_images(training_images, training_labels, training_labels, "Training Data" )
```



```
[6] display_images(validation_images, validation_labels, validation_labels, "Training Data" )
```



#### ▼ Preprocess Dataset

Here, you'll perform normalization on images in training and validation set.

- You'll use the function [preprocess\\_input](#) from the ResNet50 model in Keras.

```
[7] def preprocess_image_input(input_images):
    input_images = input_images.astype('float32')
    output_ims = tf.keras.applications.resnet50.preprocess_input(input_images)
    return output_ims
```

```
[8] train_X = preprocess_image_input(training_images)
valid_X = preprocess_image_input(validation_images)
```

#### ▼ Define the Network

You will be performing transfer learning on **ResNet50** available in Keras.

- You'll load pre-trained **imagenet weights** to the model.
- You'll choose to retain all layers of **ResNet50** along with the final classification layers.

```
[9] ...
Feature Extraction is performed by ResNet50 pretrained on imagenet weights.
Input size is 224 x 224.
...
def feature_extractor(inputs):

    feature_extractor = tf.keras.applications.resnet.ResNet50(input_shape=(224, 224, 3),
                                                               include_top=False,
                                                               weights='imagenet')(inputs)

    return feature_extractor

...
Defines final dense layers and subsequent softmax layer for classification.
...
def classifier(inputs):
    x = tf.keras.layers.GlobalAveragePooling2D()(inputs)
    x = tf.keras.layers.Flatten()(x)
    x = tf.keras.layers.Dense(1024, activation="relu")(x)
    x = tf.keras.layers.Dense(512, activation="relu")(x)
    x = tf.keras.layers.Dense(10, activation="softmax", name="classification")(x)
    return x

...
Since input image size is (32 x 32), first upsample the image by factor of (7x7) to transform it to (224 x 224)
Connect the feature extraction and "classifier" layers to build the model.
...
def final_model(inputs):

    resize = tf.keras.layers.UpSampling2D(size=(7,7))(inputs)

    resnet_feature_extractor = feature_extractor(resize)
    classification_output = classifier(resnet_feature_extractor)

    return classification_output

...
Define the model and compile it.
Use Stochastic Gradient Descent as the optimizer.
Use Sparse Categorical CrossEntropy as the loss function.
...
def define_compile_model():
    inputs = tf.keras.layers.Input(shape=(32,32,3))

    classification_output = final_model(inputs)
    model = tf.keras.Model(inputs=inputs, outputs = classification_output)

    model.compile(optimizer='SGD',
```

```

        loss='sparse_categorical_crossentropy',
        metrics = ['accuracy'])

    return model

model = define_compile_model()

model.summary()

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5
94773248/94765736 [=====] - 1s 0us/step
Model: "model"

Layer (type)          Output Shape         Param #
=====
input_1 (InputLayer)   [(None, 32, 32, 3)]      0
up_sampling2d (UpSampling2D) (None, 224, 224, 3)  0
resnet50 (Functional)  (None, 7, 7, 2048)       23587712
global_average_pooling2d (GlobalAveragePooling2D) (None, 2048)  0
flatten (Flatten)     (None, 2048)           0
dense (Dense)         (None, 1024)          2098176
dense_1 (Dense)       (None, 512)           524800
classification (Dense)(None, 10)            5130
=====
Total params: 26,215,818
Trainable params: 26,162,698
Non-trainable params: 53,120

```

## ▼ Train the model

```

✓ [10] # this will take around 20 minutes to complete
33m
EPOCHS = 4
history = model.fit(train_X, training_labels, epochs=EPOCHS, validation_data = (valid_X, validation_labels), batch_size=64)

Epoch 1/4
782/782 [=====] - 529s 628ms/step - loss: 0.4005 - accuracy: 0.8692 - val_loss: 0.2145 - val_accuracy: 0.9266
Epoch 2/4
782/782 [=====] - 494s 632ms/step - loss: 0.1039 - accuracy: 0.9664 - val_loss: 0.1739 - val_accuracy: 0.9396
Epoch 3/4
782/782 [=====] - 494s 632ms/step - loss: 0.0368 - accuracy: 0.9897 - val_loss: 0.1609 - val_accuracy: 0.9494
Epoch 4/4
782/782 [=====] - 494s 632ms/step - loss: 0.0146 - accuracy: 0.9967 - val_loss: 0.2903 - val_accuracy: 0.9297

```

## ▼ Evaluate the Model

Calculate the loss and accuracy metrics using the model's `.evaluate` function.

```

✓ [11] loss, accuracy = model.evaluate(valid_X, validation_labels, batch_size=64)
41s
157/157 [=====] - 28s 176ms/step - loss: 0.2903 - accuracy: 0.9297

```

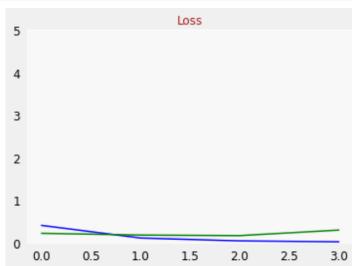
## ▼ Plot Loss and Accuracy Curves

Plot the loss (in blue) and validation loss (in green).

```

✓ [12] plot_metrics("loss", "Loss")
0s

```

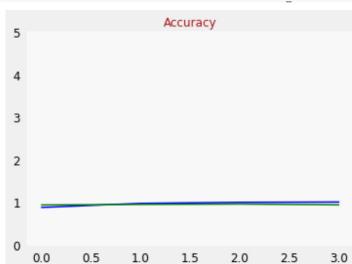


Plot the training accuracy (blue) as well as the validation accuracy (green).

```

✓ [13] plot_metrics("accuracy", "Accuracy")
0s

```



## ▼ Visualize predictions

You can take a look at the predictions on the validation set.

```
29s ✓
probabilities = model.predict(valid_X, batch_size=64)
probabilities = np.argmax(probabilities, axis = 1)

display_images(validation_images, probabilities, validation_labels, "Bad predictions indicated in red.")
```

Bad predictions indicated in red



frog



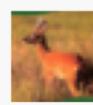
airplane



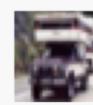
airplane



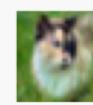
frog



deer



truck



cat



dog



deer



cat

✓ 29s completed at 9:11 PM

● X